

• EXPLORE · ATTACK · CONTAIN · REPORT

# The Agent-Native Red-Teaming Workspace.

- ▶ Explore like a hacker. Your agent maps the app through its own browser.
- ▶ Attack the traffic. It edits, replays, and compares real requests.
- ▶ Stay in bounds. Sandbox limits you set, and a record of what it did.

Professional red-teaming for everyone.

## ■ THE PROBLEM

# Professional red teaming is out of reach for most teams.

### 01 · Expertise

#### Security expertise is hard to find

Red teaming means understanding an app, forming attack hypotheses, and verifying them. Few teams have someone who does that regularly.

### 02 · Cost

#### Outside assessments are expensive

A professional pentest is valuable, but paying for one on every release is out of reach for many teams.

### 03 · Tooling

#### Existing tools were built for people

GUIs, proxies, and CA setup leave an agent reading screens and juggling tools instead of testing the app.

### 04 · Control

#### AI agents can run out of control

A mistake or a prompt injection can aim testing at the wrong systems. We need to sandbox agents.

■ WHAT H5I GIVES YOU

# Professional red-teaming for everyone.

No security expertise needed. Your own AI agent does the testing with h5i, and you review what it found.

h5i browser

## Agent-led exploration

Your agent reads pages, logs in, clicks, and fills forms to map the app and its endpoints.

h5i websec

## Direct HTTP attacks

It edits, replays, and compares the requests the browser sent, like a professional pentester.

h5i box · .h5i/env.toml

## Sandboxed execution

The agent works inside network and filesystem limits you set, and refused actions are recorded.

h5i ui · h5i project


## Intuitive report

Review findings with their evidence, then share a report written for non-specialists.

■ THE BROWSER

# Let the agent explore the app like a hacker.

THE TARGET · WHAT A PAGE LOADS



PAGE BYTES

CAPTURES EVERY REQUEST

**h5i**



OUTLINE, NOT PIXELS

THE AGENT · WHAT IT RUNS, AND WHAT COMES BACK

the agent's session

```
$ h5i browser open https://shop.test --capture
$ h5i browser snapshot
  @e1 heading "My account"
  @e7 link "Orders"
$ h5i browser click @e7 ✓
```

\* Simple-site benchmarks vs headless Chromium · see the browser design notes

**~3x**

faster reads\*

**~86%**

less peak memory\*

**@refs**

pages as outlines to act on


**Pure Rust**

no Chromium or V8

## ■ WEB SECURITY TESTING

# Let your agent red-team the web app.

The agent takes a request the app really sent, changes one field, resends it, and compares the answers, like professional hackers.



Find real vulnerabilities, **with the proof attached**. Only on apps you are **authorized to test**.

# Keep agent-led testing within bounds.

Browser policy covers browser traffic. To constrain the agent and its other tools, sandbox the whole workflow; enforcement depends on the isolation tier.

```
h5i browser open --allow
```

## Limit where the browser goes

Browser policy rejects destinations outside the target and records each denied attempt.

```
h5i box · .h5i/env.toml
```

## Isolate the whole workflow

Run the agent and its tools inside a sandbox with explicit network and filesystem limits.

```
h5i browser audit · requests
```

## Record actions and refusals

Review browser actions, HTTP traffic, and denied connections with their reasons.

```
h5i ui
```

## Review and intervene

Watch sessions live and take over sensitive browser steps when a human is needed.

■ WHY H5I

# One workflow for agent-led red teaming. Fewer tools to coordinate.

Browser automation

The agent can drive pages. HTTP testing, limits, and findings live in other tools.

Browser + proxy tools

Deep traffic inspection built around a human GUI. The agent coordinates browser, proxy, and CA setup.

Sandboxing an existing browser

Limits what the agent can reach, without HTTP inspection, replay, or findings.

**h5i**

Exploration, the HTTP attack loop, sandbox limits, and reviewable findings in one CLI your agent already knows how to run.

# Intuitive reports, clear even to non-experts.

h5i ui · project acme · Report v1

## Security Assessment Report: ACME shop

Issued v1 · 3 findings · 22 of 24 checklist items covered

### Executive summary

Any signed-in customer can read other customers' orders by changing one number in the address. This is an **IDOR**. Fix it first.

### Findings

| ID  | SEVERITY | FINDING                            | EVIDENCE  | STATUS       |
|-----|----------|------------------------------------|-----------|--------------|
| F-1 | High     | Orders readable by any customer    | E-1 · E-2 | open         |
| F-2 | Medium   | No rate limit on password reset    | E-3       | open         |
| F-3 | Low      | Session cookie missing Secure flag | E-4       | fix verified |

#### TERM

### IDOR

Changing an ID in a request reaches someone else's data.

The app trusts the order number without checking that it belongs to the person asking. The fix is an ownership check on every lookup.

[OWASP](#) ↗

✓ Proof attached to every finding

✓ Credentials removed from evidence

✓ Frozen versions · PDF export

■ TRACTION & INITIAL USERS

# Built in the open.

## Built for teams that test more often.

**650+**

GitHub stars

**60+**

forks

**20+**

contributors

**1.2k+**

downloads

GitHub · checked September 26, 2026 · [Repository and releases](#)

Built for · product teams

### No security specialist

Let their own agent test each release, then review the findings and the report.

Built for · security engineers

### Test between assessments

Run agent-led checks more often, with the evidence to verify every claim.

Built for · pentesters & researchers

### Automate the repetitive work

Hand recon and routine checks to an agent in a sandbox, for engagements and CTFs.

■ NEXT MILESTONES

# Make professional red teaming accessible to every team.

## Product

### Test more apps, more easily

Broaden website compatibility and make agent-led verification simpler to run and review.

## Adoption · planned

### Work with design partners

Goal: 10 design partners testing their own apps with h5i, leading to the first paid pilots.

## Open source

### Build with the community

Keep h5i local-first and Apache 2.0, improved alongside contributors and security researchers.

Try it with your own AI agent: `curl -fsSL https://h5i.dev/install.sh | sh -s -- --websec`

[h5i.dev](https://h5i.dev)

[github.com/h5i-dev/h5i](https://github.com/h5i-dev/h5i)